

```
PS C:\Users\Dwinayan\Desktop\All\lit_searcher> uv run python evaluate.py
USER_AGENT environment variable not set, consider setting it to identify your requests.
Generating goldens using gpt-4o-mini...
Saved 10 goldens to goldens.json
You're running DeepEval's latest Contextual Precision Metric! (using gpt-4o-mini, strict=False, async_mode=True)...
You're running DeepEval's latest Contextual Recall Metric! (using gpt-4o-mini, strict=False, async_mode=True)...
You're running DeepEval's latest Contextual Relevancy Metric! (using gpt-4o-mini, strict=False, async_mode=True)...
You're running DeepEval's latest Answer Relevancy Metric! (using gpt-4o-mini, strict=False, async_mode=True)...
You're running DeepEval's latest Faithfulness Metric! (using gpt-4o-mini, strict=False, async_mode=True)...
```

DeepEval Evaluation Results

test_case_0

Input:

Imagine if Anthropic's trademark dispute escalated, forcing a complete rebranding of AI projects. What implications would this have for innovation and identity in AI?

Actual Output:

If Anthropic's trademark dispute escalated and forced a complete rebranding of AI projects, the implications for innovation and identity in AI could be significant:

1. ****Stifling Innovation****: A requirement for widespread rebranding could hinder the momentum of AI development. Companies may become more risk-averse, fearing that their branding choices could lead to legal complications, thereby slowing the pace of innovation in the field.

2. ****Identity Confusion****: Rebranding could generate confusion among users and developers. Established brands often carry significant recognition and trust; losing a brand name can complicate user relationships and dilute the identity that companies have cultivated.

3. ****Market Fragmentation****: A wave of rebranding might lead to fragmentation within the AI market. As projects rename themselves to avoid trademark issues, users may have difficulty distinguishing between similar products, potentially leading to decreased market efficiency and increased consumer decision fatigue.

4. ****Enhanced Competition****: On the flip side, rebranding could create new opportunities for differentiation in a crowded market. Companies may take the opportunity to redefine their missions and offer unique value propositions, possibly leading to a surge in creative and innovative solutions.

5. ****Shift in Brand Awareness****: The rebranding process, especially if it is forced and publicized, could lead to increased media coverage and public discourse around the affected projects. This heightened visibility might amplify awareness and engagement with AI technologies, albeit amidst the challenges of identity transition.

6. ****Regulatory and Compliance Impact****: An escalation in trademark disputes may invite stricter scrutiny from regulatory bodies concerned with intellectual property in the rapidly evolving AI landscape. This could prompt the establishment of clearer guidelines for trademark usage, potentially fostering a more standardized approach to branding within the industry.

In summary, while an escalation of the trademark dispute could present challenges related to innovation and identity in AI, it could also catalyze market shifts and opportunities for companies to reassert their positioning in a competitive landscape.

Expected Output:

If Anthropic's trademark dispute escalated and required a complete rebranding of AI projects, it could stifle innovation by diverting resources and focus away from development towards legal and marketing efforts. Additionally, frequent rebranding could lead to confusion among users and diminish brand identity, making it harder for projects to establish a loyal user base. This situation highlights the delicate balance between protecting intellectual property and fostering a collaborative environment for AI advancement.

Metrics Status

Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	0.83	0.70	The score is 0.83 because the relevant nodes ar...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every aspect of the e...
FAIL	Contextual Relevancy	0.13	0.70	The score is 0.13 because while there are mentions of the 'Anthropic trademark dispute' and its implications for rebranding, such as 'Anthropic asked us to change our name' and 'the forced rebranding generated substantial secondary publicity', the majority of the context focuses on unrelated topics like 'societal implications around employment' and 'data privacy', which do not address the specific implications for innovation and identity in AI.
PASS	Answer Relevancy	1.00	0.70	The score is 1.00 because the response directly...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_1

Input:

What significance does the name 'Moltbot' hold regarding lobster growth, and how does it relate to AI development?

Actual Output:

I wasn't able to find relevant information in the uploaded papers to answer your question. You may want to rephrase your question or upload additional papers.

Expected Output:

The name 'Moltbot' signifies the process of molting, which is how lobsters grow by shedding their exoskeletons. This metaphorically relates to AI development as it represents transformation and evolution, paralleling the way AI systems, like Moltbot, evolve and improve over time. The name change from 'Clawdbot' to 'Moltbot' reflects this mission of growth and adaptation in the AI field.

Metrics Status

Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	1.00	0.70	The score is 1.00 because the relevant nodes ar...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every aspect of the e...
FAIL	Contextual Relevancy	0.05	0.70	The score is 0.05 because while the statement 'Molt fits perfectly – it's what lobsters do to grow' provides a direct connection to lobster growth, the majority of the retrieval context is irrelevant, focusing instead on installation and features of Moltbot without addressing its significance.
FAIL	Answer Relevancy	0.00	0.70	The score is 0.00 because the response completely fails to address the significance of 'Moltbot' in relation to lobster growth and AI development, instead providing irrelevant information that does not contribute to the question.
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_2

Input:

What are the significant findings on the project's adoption metrics, security vulnerabilities, and user experiences?

Actual Output:

The project's adoption metrics, security vulnerabilities, and user experiences reveal several significant findings:

Adoption Metrics

- ****GitHub Activity Indicators****: The report indicates a measurable increase in activity on GitHub, suggesting growing engagement and development within the community.
- ****Community Engagement****: There is active participation in Discord and Slack channels, with many community members providing support and sharing workflows. This reflects a thriving community that enhances the project's usability.
- ****Media Coverage and Mindshare****: The project has received coverage that likely contributes to its visibility and user interest, although specific details on this aspect are not provided.

Security Vulnerabilities

- ****File Integrity Monitoring and Network Traffic Analysis****: The implementation of these techniques is critical for identifying and mitigating potential security threats.
- ****Secret Management Practices****: The recommendation to utilize external secret management systems like HashiCorp Vault and AWS Secrets Manager is crucial for maintaining security and safeguarding sensitive information.
- ****Skill Auditing****: Reviewing skills code before installation helps in reducing vulnerabilities associated with untrusted code.
- ****Regular Backups****: Maintaining configuration backups is emphasized for disaster recovery, highlighting the importance of preparedness against data loss.
- ****Access Logging****: Enabling comprehensive logging and regularly reviewing logs are necessary practices to detect and respond to suspicious activities.

User Experiences

- ****User Expectations****: Non-technical users anticipate a polished consumer product experience, while organizations in regulated industries may lack dedicated security expertise. This mismatch can lead to frustrations among these groups.
- ****Security Responsibility Concerns****: Individuals and organizations that are uncomfortable with taking on security responsibilities may find the project less appealing, particularly if they are unwilling to accept risks associated with experimental software.
- ****Initial Setup and Mastery****: Users report that initial setup takes between 2 to 8 hours, depending on their technical expertise and integration needs. Effective usage requires 1-2 weeks of experimentation, indicating a learning curve that may deter some potential users.
- ****Documentation Quality****: The documentation has received mixed reviews, with some users finding it too technical and inconsistent, which can hinder onboarding.
- ****Productivity Gains vs. Friction****: Long-term users acknowledge the project offers transformative productivity capabilities despite ongoing challenges, suggesting that while the project has room for improvement, the benefits are substantial enough to retain a loyal user base.

Overall, the project shows promise in user engagement and security practices but faces challenges regarding user accessibility and experience during the onboarding process.

Expected Output: The project's adoption metrics indicate rapid growth, primarily driven by its privacy-first architecture and genuine automation capabilities. However, significant security vulnerabilities have been identified, largely due to architectural design flaws and user misconfigurations. User experiences suggest that the tool is best suited for technically sophisticated individuals who are comfortable with self-hosting and managing security, reflecting a strong developer community but an immature security posture.

Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	0.81	0.70	The score is 0.81 because while the relevant no...
PASS	Contextual Recall	0.80	0.70	The score is 0.80 because while the sentences a...
FAIL	Contextual Relevancy	0.19	0.70	The score is 0.19 because the retrieval context primarily contains irrelevant information, such as 'Implement file integrity monitoring' and section titles like '11.8 Comparative User Assessments', which do not address the requested findings on adoption metrics, security vulnerabilities, or user experiences. The only relevant statements are '10. Adoption Metrics as of January 31, 2026' and 'Long-term users consistently report transformative capability', but they are overshadowed by the lack of substantial information.
PASS	Answer Relevancy	1.00	0.70	The score is 1.00 because the response directly...
PASS	Faithfulness	0.93	0.70	The score is 0.93 because the actual output sug...

test_case_3

Input: What significant events led to the project's launch as 'Clawdbot' on November 24, 2025?
Actual Output: The report indicates that the project was originally launched as "Clawdbot" on November 24, 2025, and there were no specific significant events mentioned in the context leading up to its launch. The timeline provided focuses on events that occurred after the initial launch, including a trademark notice from Anthropic on January 27, 2026, which necessitated a rename to "Moltbot." Thus, the report does not elaborate on events or developments prior to November 24, 2025, that contributed to the launch of "Clawdbot."
Expected Output: The project was launched as "Clawdbot" on November 24, 2025, with a focus on privacy-first architecture and automation capabilities. The name, inspired by a lobster mascot named "Clawd," symbolized the tool's ability to "grab" and manipulate data and systems.

Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	1.00	0.70	The score is 1.00 because all relevant nodes ar...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every sentence in the...
FAIL	Contextual Relevancy	0.10	0.70	The score is 0.10 because while the retrieval context mentions the launch date and name of 'Clawdbot', such as 'The project was originally launched as 'Clawdbot' on November 24, 2025', it lacks details on the significant events leading to that launch, as highlighted by the numerous irrelevant statements discussing post-launch issues and name changes.
FAIL	Answer Relevancy	0.20	0.70	The score is 0.20 because the output fails to address the input question about significant events leading to the launch of 'Clawdbot,' instead focusing on irrelevant post-launch details and explicitly stating a lack of elaboration on prior events.
FAIL	Faithfulness	0.60	0.70	The score is 0.60 because the actual output implies a lack of events prior to the launch of 'Clawdbot', while the context suggests there were significant occurrences, such as the trademark notice from Anthropic, indicating notable events did happen.

test_case_4

Input: What prompted the rebranding of Clawdbot to Moltbot amid account hijacking issues?
Actual Output: The rebranding of Clawdbot to Moltbot was prompted by a trademark dispute, specifically a cease and desist order from Anthropic. Following this legal pressure, the Clawdbot team was forced to change their branding, which inadvertently created a vulnerability during the transition period. This brief "namespace vacuum" allowed opportunistic attackers to hijack the original social media accounts and repositories associated with Clawdbot, leading to significant security issues, including scams being perpetrated under the hijacked accounts.
Expected Output: The rebranding of Clawdbot to Moltbot was prompted by account hijacking chaos, where crypto scammers took over the old social media handles to promote fraudulent schemes. This situation necessitated a new identity to mitigate confusion and protect the platform's integrity.

Metrics Status	Metric	Score	Threshold	Reason
----------------	--------	-------	-----------	--------

PASS	Contextual Precision	0.80	0.70	The score is 0.80 because while the relevant no...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because the expected output d...
FAIL	Contextual Relevancy	0.46	0.70	The score is 0.46 because while some relevant statements like 'Clawdbot was forced to rename to Moltbot due to a trademark dispute' and 'One legal notice can force a rebrand that exposes you to account hijacking' provide context for the rebranding, the majority of the retrieval context focuses on cybersecurity issues and unrelated details, such as 'The code is clean. The infrastructure is not.' which do not address the reasons for the rebranding.
PASS	Answer Relevancy	0.75	0.70	The score is 0.75 because while the response ad...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_5				
Input:				
Actual Output:				
Imagine if Moltbot could autonomously debug complex software issues in real-time. What tasks might it perform? If Moltbot could autonomously debug complex software issues in real-time, it might perform the following tasks:				
1. **Automatic Detection of Bugs** : Utilizing Sentry webhooks to detect production bugs automatically without human intervention.				
2. **Code Analysis** : Analyzing the affected code to determine the root cause of the issue, potentially using integrated AI models for deeper insights.				
3. **Autonomous Code Fixing** : Implementing code fixes based on analysis results, allowing the system to correct issues without manual coding intervention.				
4. **Progress Updates** : Sending real-time progress updates to developers through messaging platforms, such as Telegram, to keep them informed of the debugging status.				
5. **Pull Request Creation** : Creating pull requests with proposed fixes already integrated, streamlining the process for easy review and deployment.				
6. **Execution of Tests** : Running predefined test suites to validate that code changes resolve the issues and do not introduce new bugs.				
7. **Error Reporting and Logging** : Maintaining logs of all detected issues, actions taken, and outcomes to create a historical database for future reference.				
8. **Integration with Development Tools** : Coordinating with version control systems and development environments to streamline debugging processes, ensuring developers can review and deploy fixes efficiently.				
9. **Session Management** : Maintaining isolated sessions for debugging different projects or components simultaneously, providing context and history for each debugging attempt.				
Overall, these tasks would improve the efficiency of software debugging processes, reduce downtime, and allow developers to focus on higher-level tasks while Moltbot manages real-time debugging autonomously.				
Expected Output: If Moltbot could autonomously debug complex software issues in real-time, it might perform tasks such as executing shell commands to identify errors, running and debugging code with iterative refinement to fix bugs, and automating the testing of software to ensure functionality. Additionally, it could analyze logs for anomalies and suggest optimizations based on its findings.				
Metrics				
Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	1.00	0.70	The score is 1.00 because all relevant nodes ar...
FAIL	Contextual Recall	0.17	0.70	The score is 0.17 because while there is a connection to analyzing logs for anomalies in node 4, the majority of the expected output discusses concepts like autonomous debugging and software testing that are not present in the retrieval context.
FAIL	Contextual Relevancy	0.22	0.70	The score is 0.22 because the majority of the retrieval context is irrelevant, focusing on tasks like 'unsubscribing, archiving, and prioritization' which do not relate to debugging. However, there are relevant statements such as 'Development Workflows: Developers trigger autonomous code fixing from mobile devices' and 'Production Bug Fixing: Automatic detection via Sentry webhooks' that hint at debugging capabilities, but they are overshadowed by the unrelated content.
PASS	Answer Relevancy	1.00	0.70	The score is 1.00 because the response directly...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_6				
Input:				
Actual Output:				
Imagine if the report date shifted to a year later. How would that affect the assessment's relevance? If the report date shifted to a year later, to January 31, 2027, several aspects of the assessment's relevance could be affected:				
1. **Technological Advancements** : The AI landscape is rapidly evolving, and new developments in AI technologies, frameworks, and best practices may arise within that year. The analysis provided in the report might become outdated if newer tools or methodologies emerge that outperform or enhance the capabilities of Moltbot.				
2. **User Feedback Evolution** : Long-term user experiences and feedback could change over time as users continue to interact with Moltbot. New user reviews may reflect different sentiments, especially regarding ongoing issues such as code quality and complexity. This could lead to a shift in both positive and negative assessments, altering the overall perception of the product.				
3. **Security Concerns** : Given the mention of substantial security concerns and deployment complexities, any incidents or advancements in security practices that take place over the additional year could significantly impact the assessment's relevance. New vulnerabilities or improvements in security protocols may affect user trust and adoption rates.				
4. **Competitive Landscape** : The competitive context may shift dramatically as newer open-source projects or proprietary tools emerge. If competitors introduce superior features or ease of use, Moltbot's standing in the market could change, impacting its adoption metrics and long-term viability.				
5. **Adoption Metrics** : Actual adoption rates over the additional year will provide clearer insights into Moltbot's growth trajectory. Changes in user numbers, engagement levels, and the emergence of enterprise pilots could provide a more accurate picture of its market position and user satisfaction.				
6. **Research and Development Outcomes** : Ongoing enhancements to Moltbot's features based on user feedback and technological progress could either bolster or hinder the assessments made in the initial report. Any improvements or failures in implementing desired features would affect user experiences and the product's perceived value.				
In summary, shifting the report date to a year later would likely require a reevaluation of all key findings, addressing				

Expected Output:		new developments, evolving user perspectives, changing market dynamics, and the implications of ongoing research and development efforts. Shifting the report date to January 31, 2027, could impact the assessment's relevance by introducing new developments in the rapidly evolving AI landscape. Given Moltbot's significant growth and the ongoing security concerns and trademark disputes, a later date would allow for updated insights into its capabilities, user adoption, and any emerging risks or competitive dynamics. This could enhance the report's value by reflecting the most current state of the project and its ecosystem.		
Metrics Status	Metric	Score	Threshold	Reason
FAIL	Contextual Precision	0.50	0.70	The score is 0.50 because while the second node is relevant and discusses the implications of shifting the report date, the first, third, fourth, and fifth nodes are all irrelevant and ranked lower. Specifically, the first node states, "The first document does not provide any relevant information regarding the assessment's relevance or the impact of shifting the report date," which indicates a lack of pertinent content. Similarly, the third node mentions user experiences, the fourth discusses future developments without relevance, and the fifth focuses on user feedback, all of which do not address the core question. This mix of relevant and irrelevant nodes leads to a moderate score.
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every aspect of the e...
FAIL	Contextual Relevancy	0.03	0.70	The score is 0.03 because the retrieval context primarily contains irrelevant information, such as 'updating dependencies' and 'user experiences', which do not address the impact of a shifted report date on assessment relevance. The only relevant statement is 'Report Date: January 31, 2026', but it lacks context to connect it to the assessment's relevance.
PASS	Answer Relevancy	1.00	0.70	The score is 1.00 because the response directly...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_7		Input: Actual Output:		
		Imagine if Moltbot's features were replicated by competitors overnight. How would this affect its GitHub popularity? If competitors were to replicate Moltbot's features overnight, it could significantly impact its GitHub popularity in several ways:		
		1. **Increased Competition** : The sudden availability of similar capabilities from competitors could dilute Moltbot's unique selling proposition. Developers might explore alternatives that offer the same functionalities, diminishing the sense of exclusivity that Moltbot currently enjoys.		
		2. **Shifts in Developer Interest** : Given that Moltbot's growth was fueled by a narrative of empowerment and impressive productivity, the novelty and curiosity factor might decline as other options emerge. This could lead to a decrease in new users and contributions, as developers may gravitate towards newer or differentiated offerings that promise similar or better value.		
		3. **Community Engagement** : Moltbot's viral momentum benefited from community discussions and extensive media coverage. If competitors emerge with similar tools, developer engagement might get fragmented across multiple options, reducing the community's focus on Moltbot. This could lead to fewer forks, commits, and stars as the developer community diversifies its interests.		
		4. **Sustained Virality Challenges** : Moltbot's rapid growth metrics, such as reaching over 103,000 stars in three months, rely on maintaining a unique edge. If competitors replicate its features but also provide compelling improvements or enhancements, Moltbot might struggle to sustain its "viral" status in GitHub conversations and trends.		
		5. **Innovation Pressure** : Moltbot would face increased pressure to innovate and differentiate itself continually. If it cannot adapt quickly or enhance its features beyond what has been replicated, it might see a plateau or decline in popularity.		
Expected Output:		Overall, while Moltbot has achieved significant momentum, the emergence of similar offerings by competitors could lead to a decrease in its GitHub popularity unless it continues to innovate and leverage its unique community and productivity narrative. If Moltbot's features were replicated by competitors overnight, its GitHub popularity could potentially decline. The unique capabilities that have driven its viral growth, such as proactive automation and multi-channel messaging integration, would lose their competitive edge. This could lead to a decrease in user interest and engagement, impacting the number of stars and overall adoption metrics. Additionally, the presence of alternatives may shift developer focus away from Moltbot, especially if those competitors address existing security concerns and deployment complexities more effectively.		
Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	0.95	0.70	The score is 0.95 because the relevant nodes ar...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every aspect of the e...
FAIL	Contextual Relevancy	0.70	0.70	The score is 0.70 because while there are relevant statements about Moltbot's exceptional growth and popularity on GitHub, such as 'Moltbot achieved exceptional growth velocity that ranks among the fastest in GitHub history,' many statements in the retrieval context do not directly address the impact of competitors replicating its features, leading to a mixed relevance.
PASS	Answer Relevancy	1.00	0.70	The score is 1.00 because there are no irreleva...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_8		Input: Actual Output:		
		Imagine if the project had never rebranded from Clawdbot to Moltbot. How would its identity have evolved? If the project had never rebranded from Clawdbot to Moltbot, its identity would likely have evolved differently due to the following factors highlighted in the report:		
		1. **Brand Continuity** : Clawdbot would have maintained its original brand identity, which may have strengthened recognition among early users and allowed for a more gradual evolution of its capabilities and offerings. The brand would continue to be associated with its initial launch and services without the disruption caused by the rebranding.		
		2. **Avoidance of Confusion** : The rapid transition from Clawdbot to Moltbot and then to OpenClaw resulted in significant community confusion. Without the rebranding, users would not have faced the challenge of keeping up with new names, documentation updates, and potential mismatches in resources, which could foster greater user trust and loyalty.		
		3. **Response to Trademark Issues** : Remaining as Clawdbot might have prompted a different response to the trademark notice from Anthropic. The project could have engaged in legal negotiations or alterations to its branding strategy that still aligned with the original name, rather than a complete name change.		
		4. **Market Positioning and Messaging** : The identity of Clawdbot would likely have evolved based on its core functionalities and market demands without the need for symbolic restructuring represented by the rebranding. The		

messaging could have continued to focus on the personal assistant aspects powered by Anthropic's Claude API, without the distractions of identity crises.

5. ****User Experience****: The transition challenges, such as account hijacking incidents and disruptions during the migration, could have been avoided or mitigated. A stable identity might have led to a more seamless user experience, enhancing customer satisfaction and potentially expanding the user base organically.

6. ****Development Trajectory****: Clawdbot's development trajectory might have been more focused on enhancing capabilities associated with the original branding rather than adapting to new brand identities. This could have led to a more cohesive feature set and user interface.

In summary, if Clawdbot had not rebranded, it might have capitalized on the initial launch momentum, potentially offering a more stable and consistent growth path with fewer disruptions and greater user trust. This could have resulted in a stronger community and a more focused development approach.

Expected Output:

If the project had never rebranded from Clawdbot to Moltbot, its identity would likely have evolved around the original name, maintaining a consistent brand image. The absence of the rapid rebranding could have fostered a more stable community and user base, potentially avoiding the chaos associated with the account hijacking incident. However, it may have faced challenges related to the trademark notice from Anthropic, which prompted the initial name change. Overall, the identity would have been less fragmented, but it might have struggled with legal issues under the Clawdbot name.

Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	1.00	0.70	The score is 1.00 because all relevant nodes ar...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every aspect of the e...
FAIL	Contextual Relevancy	0.52	0.70	The score is 0.52 because while there are relevant statements about the rebranding process, such as 'The rebranding from Clawdbot to Moltbot occurred rapidly within hours on January 27, 2026.' and 'The new name 'Moltbot' references the biological process of molting—when lobsters shed their shells to grow—symbolizing the project's transformation.', the majority of the context focuses on unrelated issues like account hijacking and technical capabilities, which detracts from addressing the evolution of the project's identity.
PASS	Answer Relevancy	1.00	0.70	The score is 1.00 because the response directly...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

test_case_9

Input:

What prompted the rapid rebranding from Clawdbot to Moltbot and then to OpenClaw, and what chaos ensued?

Actual Output:

The rapid rebranding from Clawdbot to Moltbot was prompted by a trademark notice from Anthropic, the company behind Claude AI, which requested a name change due to the similarity between "Clawd" and "Claude." The project creator, Peter Steinberger, chose the name "Moltbot" to symbolize transformation, likening it to the biological process of molting where lobsters shed their shells to grow.

The transition to Moltbot created significant chaos, including:

- **Account hijacking****: Scammers took control of the old "Clawdbot" social media handles on X (Twitter) and GitHub, promoting cryptocurrency schemes.
- **GitHub account issues****: The migration caused temporary disruptions in Steinberger's personal GitHub account.
- **Community confusion****: The rebranding necessitated rapid updates to documentation, tutorials, and resources to reflect the new name.
- **Domain migration****: The project's domain shifted from clawd.bot to molt.bot, requiring careful coordination.

Just three days after becoming Moltbot, the project underwent another rebranding to "OpenClaw," indicating ongoing naming challenges or strategic repositioning, although its core functionality and codebase remained unchanged.

Expected Output:

The rapid rebranding from Clawdbot to Moltbot was prompted by an Anthropic trademark notice received on January 27, 2026, which forced the first rename. The transition to Moltbot created significant chaos, including account hijacking issues. Just three days later, on January 30, 2026, the project underwent another rebranding to OpenClaw, indicating ongoing naming challenges despite the core functionality remaining unchanged.

Metrics Status	Metric	Score	Threshold	Reason
PASS	Contextual Precision	1.00	0.70	The score is 1.00 because all nodes in the retr...
PASS	Contextual Recall	1.00	0.70	The score is 1.00 because every sentence in the...
FAIL	Contextual Relevancy	0.61	0.70	The score is 0.61 because while there are relevant statements such as 'The rebranding from Clawdbot to Moltbot occurred rapidly within hours on January 27, 2026.' and 'The transition created significant chaos.', the majority of the retrieval context focuses on tool capabilities and unrelated information, making it less relevant to the specific inquiry about the reasons for rebranding and the ensuing chaos.
PASS	Answer Relevancy	0.92	0.70	The score is 0.92 because while the response pr...
PASS	Faithfulness	1.00	0.70	The score is 1.00 because there are no contradi...

Aggregate Metrics

Metric	Average Score	Pass Rate	Total
Contextual Precision	0.89	90.00% passed=9 failed=1	10
Contextual Recall	0.90	90.00% passed=9 failed=1	10
Contextual Relevancy	0.30	0.00% passed=0 failed=10	10
Answer Relevancy	0.79	80.00% passed=8 failed=2	10
Faithfulness	0.95	90.00% passed=9 failed=1	10

⚠ WARNING: No hyperparameters logged.

» Log hyperparameters to attribute prompts and models to your test runs.

✓ Evaluation completed 🕒 (time taken: 54.69s | token cost: 0.04415670000000001 USD)

» Test Results (10 total tests):

» Pass Rate: 0.0% | Passed: 0 | Failed: 10

» Want to share evals with your team, or a place for your test cases to live? ♥️ 📄

» Run 'deepeval view' to analyze and save testing results on Confident AI.

Results saved to eval_results.json.